

1. Design suitable Data structures and implement Pass-I of two pass assembler to generate intermediate code. Sample input is following assembly language program.
START 200 MOVER AREG, ='5' MOVEM AREG, X L1 MOVER BREG, ='2' ORIGIN L1+3
LTORG NEXT ADD AREG, ='1' SUB BREG, ='2' BC LT, BACK LTORG BACK EQU L1 ORIGIN
NEXT+5 MULT CREG, ='4' STOP X DS 1 END

Aim:

To implement **Pass-I of Two Pass Assembler** to generate **Intermediate Code, Symbol Table and Literal Table**.

Theory (Say in viva in 20 seconds):

Pass-I of the assembler **scans the assembly program line-by-line**, assigns addresses using **Location Counter (LC)**, fills the **Symbol Table**, collects **Literals** in the **Literal Table**, and generates **Intermediate Code (IC)**.

No machine code is generated in this phase.

Literal addresses are assigned at **LTORG / END**.

Algorithm:

1. Initialize **LC = starting address** specified in START instruction.
2. For each line in source program:
 - o If label exists → enter into **Symbol Table** with current LC.
 - o Identify instruction type (Imperative, Declarative, Assembler Directive).
 - o Generate **Intermediate Code** accordingly.
 - o If literal appears (=value) → store it in **Literal Table**.
3. On encountering **LTORG or END**, assign addresses to unassigned literals.
4. Print **SYMTAB, LITTAB, and IC**.

2. Design Pass II of a two pass assembler for following intermediate code and data structures to generate machine code. Symbol table LITERAL TABLE Symbol address LITERAL ADDRESS X 214 ='5' 205 L1 202 ='2' 206 NEXT 207 ='1' 210 BACK 202 ='2' 211 ='4' 215 Intermediate code (AD,01) (C,200) (IS,04) 1 (L,1) (IS,05) 1 (S,1) (IS,04) 2(L,2) (AD,03) (S,2)+3 (AD,05) (L,1) (L,2) (IS,01) 1 (L,3) (IS,02) 2 (L,4) (IS,07) 1(S,4) (AD,05) (L,3) (L,4) (AD,04) (S,2) (IS,03) 3 (L,5) (IS,00) (DL,02) (C,1) (AD,02)

Aim

To design and implement **Pass-II of the two-pass assembler** that reads the **Intermediate Code (IC)** produced in Pass-I and generates the **final machine code** using **Symbol Table** and **Literal Table**.

Theory

Pass-II converts **mnemonics and symbolic addresses** into their **actual machine code values**. It reads:

- **Intermediate Code (IC)**
- **Opcode Table (OPTAB)**
- **Symbol Table (SYMTAB)**
- **Literal Table (LITTAB)**

Pass-II replaces:

(IS, opcode) → Actual machine opcode

(S, index) → Address from Symbol Table

(L, index) → Address from Literal Table

(C, value) → Constant value

Assembler directives such as **START, END, ORIGIN, EQU, LTORG** do **not** produce machine code.

Algorithm (Simple and Viva-Friendly)

1. Read **Intermediate Code line by line**.
2. For each line:
 - Identify instruction class:

- IS → **Imperative Instruction**
- DL → **Declarative Instruction**
- AD → **Assembler Directive**

3. If the line is IS:

- Convert opcode → machine opcode
- Convert register → register code
- Replace symbol/literal reference → actual address

4. If DL:

- Allocate memory but do not generate machine code

5. Print final **Machine Code** instruction line-by-line.

Practical 3 – Pass-I of Macro Processor

Aim

To design and implement **Pass-I of a two-pass macro processor** to construct:

- **MNT (Macro Name Table)**
 - **MDT (Macro Definition Table)**
 - **ALA (Argument List Array)**
-

Theory

A **macro** is a sequence of assembly instructions that is given a name and can be reused anywhere in the program.

Macro Processor **expands** these macros before actual assembly.

The **Macro Processor** works in **two passes**:

Pass Purpose

Pass-I Identify macro definitions and store them in **MNT** and **MDT**, convert parameters into placeholders.

Pass-II Perform **macro expansion** by replacing macro invocation with corresponding instructions from MDT.

Algorithm (Pass-I)

1. Read source program **line-by-line**.
2. When encountering MACRO:
 - o Read next line → macro header.
 - o Extract **macro name** and its parameters.
 - o Enter macro name into **MNT** with current **MDT index**.
 - o Store parameters in **ALA** as positional #1, #2, #3
3. Copy each instruction of macro body into **MDT** replacing &ARG with corresponding #i.
4. When MEND is encountered:
 - o End definition and write MEND into MDT.
5. Continue scanning until end of file.

Practical 4 – Pass-II of Macro Processor

Aim

To perform **macro expansion** using the tables generated in Pass-I (MNT, MDT, ALA) and produce the **final expanded source program**.

Theory

In **Pass-II**, the macro processor:

- Reads the input program again
- Whenever a **macro call** is encountered, it:
 1. **Looks up** the macro name in **MNT** to find where its definition starts in **MDT**
 2. Uses **ALA** to map **actual parameters** to **formal parameters**
 3. Copies macro body from **MDT** into output, replacing **positional parameters (#1, #2, ...)** with actual arguments
- Non-macro lines are copied to output **directly** without changes.

No new tables are created here.

Pass-II only **expands** macros to produce the final code.

Algorithm (Pass-II)

1. Read program line-by-line.
2. If the line is **not** a macro call:
 - Copy it to output as is.
3. If the line **is** a macro call:
 - Retrieve its entry in **MNT** to get starting index of body in **MDT**.
 - Fill **ALA** with actual arguments.
 - For each line in **MDT** until **MEND**:
 - Replace positional parameters (#1, #2, ...) with actual arguments.
 - Write the expanded line to output.
4. Continue until **END** is reached.

Practical 5 – Synchronization Using Mutex

Aim

To solve the classical synchronization problem by using a **Mutex (Mutual Exclusion lock)** to ensure that only one process accesses the critical section at a time.

Theory

A **mutex** is a binary locking mechanism used to prevent **race conditions**.

When a process enters its **critical section**, it locks the mutex so that **no other process** can access shared data at the same time.

After finishing, the process unlocks the mutex, allowing the next waiting process to enter.

Algorithm

1. Initialize **mutex = 1**.
2. When a process wants to enter the **critical section**, it checks:
 - o If **mutex == 1**, lock it (mutex = 0).
3. Execute the **critical section** safely.
4. After completing, **release** the mutex (mutex = 1).
5. Other processes wait until mutex becomes free.
6. Continue until all processes finish execution.

Practical 6 – Synchronization Using Semaphore

Aim

To solve the classical synchronization problem using **Semaphores** to coordinate multiple processes accessing shared resources.

Theory

A **semaphore** is an integer variable used for **process synchronization**.

It controls access to shared resources by using two atomic operations:

- **wait() / P()** → decreases semaphore value → blocks if value < 0
- **signal() / V()** → increases semaphore value → wakes up blocked process

Semaphores prevent **race conditions** and ensure **mutual exclusion** in the critical section.

Algorithm

1. Initialize semaphores:

- **mutex = 1** (for locking critical section)
- **full = 0** (items produced)
- **empty = n** (space available)

2. Producer Process:

- **wait(empty)**
- **wait(mutex)**
- **produce item (critical section)**
- **signal(mutex)**
- **signal(full)**

3. Consumer Process:

- **wait(full)**
- **wait(mutex)**
- **consume item (critical section)**
- **signal(mutex)**
- **signal(empty)**

4. Repeat steps until all items are produced and consumed.

Practical 7 – CPU Scheduling: FCFS (First Come First Served)

Aim

To simulate the **First Come First Served (FCFS)** CPU scheduling algorithm and calculate **Waiting Time** and **Turnaround Time** for each process.

Theory

In **FCFS scheduling**, processes are executed in the **order of their arrival** in the ready queue. It is **non-preemptive**, meaning once a process gets the CPU, it runs until completion. This scheduling can lead to the **Convoy Effect**, where shorter processes wait behind longer ones.

Algorithm

1. Sort processes in **increasing order of Arrival Time**.
2. Initialize **time = 0**.
3. For each process:
 - o If the CPU is idle and the process has not yet arrived, set time = arrival_time.
 - o Compute **Waiting Time**:
 - o $WT = time - arrival_time$
 - o Compute **Completion Time**:
 - o $CT = time + Burst_Time$
 - o Compute **Turnaround Time**:
 - o $TAT = CT - arrival_time$
 - o Update **time = CT**.
4. Display **WT** and **TAT** for all processes.

Practical 8 – CPU Scheduling: SJF (Preemptive) / SRTF (Shortest Remaining Time First)

Aim

To simulate the **Shortest Job First (Preemptive)** CPU scheduling algorithm and calculate **Waiting Time** and **Turnaround Time** for each process.

Theory

In **SJF Preemptive**, also called **SRTF**, at any given time the **process with the shortest remaining burst time** is selected to execute.

If a new process arrives with a **shorter burst time** than the currently running one, the CPU is **preempted** and assigned to that new process.

This algorithm minimizes **average waiting time**, but requires knowledge of burst times in advance.

Algorithm

1. Maintain arrays for **Arrival Time (AT)** and **Remaining Time (RT)** for all processes.
2. Set **current time = 0, completed = 0**.
3. At each time unit:
 - o From all processes that **have arrived** and **not completed**, select the process with **minimum remaining time**.
 - o Decrease its remaining time by 1 (process executes for 1 unit).
4. If a process's remaining time becomes 0:
 - o Mark process as **completed**.
 - o Calculate:
 - o Completion Time (CT) = current time
 - o Turnaround Time (TAT) = CT - AT
 - o Waiting Time (WT) = TAT - Burst Time
5. Repeat until all processes are completed.
6. Display **WT** and **TAT**.

Practical 9 – CPU Scheduling: Priority (Non-Preemptive)

Aim

To simulate the **Priority Scheduling (Non-Preemptive)** CPU scheduling algorithm and calculate **Waiting Time** and **Turnaround Time** for each process.

Theory

In **Priority Scheduling**, each process is assigned a **priority value**.

The CPU is allocated to the process with the **highest priority** (usually **lowest priority number = highest priority**).

This algorithm is **non-preemptive**, meaning once a process starts execution, it runs until completion.

If multiple processes have the same priority, the one that arrived earlier is selected.

Algorithm

1. Accept input of processes with:
 - o Process ID
 - o Arrival Time
 - o Burst Time
 - o Priority Value
2. Sort all processes by:
 - o **Arrival Time**, and if equal,
 - o **Priority (lowest number first)**.
3. Set **current time = 0**.
4. For each process in sorted order:
 - o If current time < arrival time, set current time = arrival time.
 - o Compute:
 - o Waiting Time (WT) = current time - Arrival Time
 - o Completion Time (CT) = current time + Burst Time
 - o Turnaround Time (TAT) = CT - Arrival Time
 - o Update current time = CT.
5. Display **WT** and **TAT** for each process.

Practical 10 – CPU Scheduling: Round Robin (Preemptive)

Aim

To simulate the **Round Robin** CPU scheduling algorithm using a given **time quantum**, and calculate **Waiting Time** and **Turnaround Time** for each process.

Theory

Round Robin scheduling assigns each process a fixed **time quantum**.

Processes are executed in a **circular queue**.

If a process does not finish within its allotted time slice, it is **preempted** and placed back into the ready queue.

This ensures **fair CPU sharing** among all processes and is widely used in **time-sharing systems**.

Algorithm

1. Accept input of:
 - o Process ID
 - o Arrival Time
 - o Burst Time
 - o Time Quantum (Q)
2. Initialize:
 - o Remaining Time = Burst Time of each process
 - o Current Time = 0
 - o Ready Queue = empty
3. Add processes to the ready queue as their **arrival time ≤ current time**.
4. While there are unfinished processes:
 - o Take the **first process** from the ready queue.
 - o Execute it for:
 - o $\min(\text{Time Quantum}, \text{Remaining Time})$
 - o Reduce Remaining Time.
 - o Increase Current Time accordingly.
 - o If Remaining Time becomes 0:
 - Compute:

- Completion Time (CT) = current time
 - Turnaround Time (TAT) = CT - AT
 - Waiting Time (WT) = TAT - BT
- Else:
 - Place the process back into the queue.
 - Add any newly arrived processes to the queue.
5. Display **WT** and **TAT** for all processes.

Practical 11 – Memory Allocation: Best Fit

Aim

To simulate the **Best Fit** memory allocation strategy and allocate processes to memory blocks with **minimum internal fragmentation**.

Theory

In **Best Fit**, each process is allocated to the **smallest memory block** that is **large enough** to accommodate it.

This helps in reducing **internal fragmentation**, because we try to fit the process in the block that leaves the **least leftover space**.

However, Best Fit may lead to **external fragmentation**, since many small unused gaps can be left across memory.

Algorithm

1. Accept sizes of memory blocks.
2. Accept sizes of processes to be allocated.
3. For each process:
 - o Search all memory blocks.
 - o Select the block that:
 - Is **greater than or equal** to the process size.
 - Has the **smallest size** among all eligible blocks.
4. If such a block exists:
 - o Allocate the process to that block.
 - o Reduce the block size by the allocated process size.
5. If no suitable block is found:
 - o The process **cannot be allocated**.
6. Display allocation results.

Practical 12 – Memory Allocation: First Fit

Aim

To simulate the **First Fit** memory allocation strategy and allocate processes to memory blocks in the order they appear.

Theory

In **First Fit**, memory blocks are scanned **from the beginning**, and a process is allocated to the **first available block** that is **large enough** to hold it.

It is **simple** and **fast** because it does not search for the optimal block — it stops as soon as it finds a suitable one.

However, it can lead to **external fragmentation**, as memory becomes scattered with small free gaps.

Algorithm

1. Take input for sizes of memory blocks.
2. Take input for sizes of processes to be allocated.
3. For each process:
 - o Scan the list of memory blocks from the start.
 - o If a block is **large enough**:
 - Allocate the process to that block.
 - Reduce block size by process size.
 - Stop searching for this process.
4. If no suitable block is found:
 - o The process **remains unallocated**.
5. Display allocation results.

Practical 13 – Memory Allocation: Next Fit

Aim

To simulate the **Next Fit** memory allocation strategy and allocate processes to memory blocks starting the search from the **last allocated position**.

Theory

Next Fit works similar to **First Fit**, but instead of always starting the search from the **beginning** of the memory, it remembers the **last block** where a process was allocated. For the next process, the search begins **from that point onward**, looping back to the start if needed.

This strategy helps in reducing time spent repeatedly scanning memory from the beginning, but may still lead to **external fragmentation**.

Algorithm

1. Take input for sizes of memory blocks.
2. Take input for sizes of processes to be allocated.
3. Maintain an index **j** that marks the **last allocated block position**.
4. For each process:
 - o Starting from block index **j**, scan memory sequentially.
 - o If the block is **large enough**:
 - Allocate the process to that block.
 - Reduce the block size by process size.
 - Update **j** to the current block position.
 - Stop searching for this process.
 - o If the end of the list is reached, wrap around to the start.
5. If no suitable block is found, the process remains **unallocated**.
6. Display allocation results.

Practical 14 – Memory Allocation: Worst Fit

Aim

To simulate the **Worst Fit** memory allocation strategy and allocate processes to the **largest available memory block**.

Theory

In **Worst Fit**, each process is allocated to the **largest block** that is big enough to hold it. This aims to **reduce external fragmentation** by leaving larger remaining spaces available for future allocations rather than creating many small unusable gaps.

However, this method may still produce fragmentation but tends to create **larger free blocks**.

Algorithm

1. Input the sizes of memory blocks.
2. Input the sizes of processes.
3. For each process:
 - o Search all memory blocks.
 - o Identify the block that:
 - Is **greater than or equal** to the process size.
 - Has the **largest size** among all such blocks.
4. If such a block exists:
 - o Allocate the process to that block.
 - o Reduce the block size by the process size.
5. If no suitable block is found:
 - o Mark the process as **Not Allocated**.
6. Display allocation results.

Practical 15 – Page Replacement: FIFO (First In First Out)

Aim

To simulate the **FIFO Page Replacement** algorithm and count the number of **page faults**.

Theory

In **FIFO**, the page that has been in memory for the **longest time** is the one that gets **replaced first** when a new page needs to be loaded.

A **queue** is used to track the order of pages.

FIFO is simple but can cause **Belady's Anomaly**, where increasing the number of frames can **increase** page faults.

Algorithm

1. Take the **reference string** and number of **frames**.
2. Initialize an empty **queue** to store pages currently in memory.
3. For each page in the reference string:
 - o If the page is **already in memory**, do nothing.
 - o Else:
 - If memory is **not full**, add the page to the queue.
 - If memory is **full**, remove the **front (oldest)** page and insert the new page.
 - o Count page faults whenever a new page is added.
4. Display total **page faults**.

Practical 16 – Page Replacement: Optimal

Aim

To simulate the **Optimal Page Replacement** algorithm and count the number of **page faults**.

Theory

The **Optimal Page Replacement** algorithm replaces the page that will **not be used for the longest duration in the future**.

It gives the **minimum possible number of page faults**, meaning **no other algorithm can perform better** for a given reference string.

However, it is **not practical** in real systems because it requires **future knowledge** of memory references.

Algorithm

1. Take the **reference string** and number of **frames**.
 2. Initialize memory frames as **empty**.
 3. For each page in the reference string:
 - o If the page is **already in memory**, continue.
 - o Else:
 - If there is a **free frame**, insert the page.
 - Otherwise:
 - For each page currently in memory, look **ahead** in the reference string to determine when it will be used next.
 - Replace the page that will be used **farthest in the future** (or **not used again**).
 - o Increment page fault count whenever a page is replaced or newly inserted.
4. Display total **page faults**.

Practical 17 – Page Replacement: LRU (Least Recently Used)

Aim

To simulate the **LRU (Least Recently Used)** page replacement algorithm and count the number of **page faults**.

Theory

In **LRU**, the page that has **not been used for the longest time in the past** is replaced first. The assumption is that pages used recently are **more likely to be used again**, so LRU tries to minimize page faults by replacing the **least recently accessed** page.

It requires tracking the **order of page usage**, often maintained using a **stack or list**.

Algorithm

1. Take the **reference string** and number of **frames**.
2. Initialize the set or list of pages in memory as **empty**.
3. For each page in the reference string:
 - o If the page is **already in memory**:
 - Update its position to mark it as **most recently used**.
 - o Else:
 - If memory has a **free frame**, insert the page.
 - If memory is **full**, remove the **least recently used** page.
 - Insert the new page.
 - Increment the **page fault counter**.
4. Continue until all references are processed.
5. Display the total **page faults**.